

CME 213

SPRING 2019

Eric Darve

So far in CME213

- Multi-threaded programs: based on creating threads that run user-provided functions
- C++ threads:

```
thread t1(f1);
```

- Pthreads:

```
pthread_create(&thread[t], NULL, PrintHello, (void *)t);
```

Pthreads basics

- Include the header file:

```
include <pthread.h>
```

- Compile using:

```
gcc -o hello_pthread hello_pthread.c -lpthread
```

```
hello_pthread.c
```

```
for (t = 0; t < n_thread; t++)
{
    printf("In main: creating thread %ld\n", t);
    pthread_create(&thread[t], NULL, PrintHello, (void *)t);
}

for (t = 0; t < n_thread; t++)
{
    pthread_join(thread[t], (void **)(&thread_result) /*NULL is common*/);
    printf("Thread # %ld just finished; its value is %ld\n", t,
           thread_result);
    assert(ComplexCalculation(t) == thread_result); /* Testing output */
}
```

Reference: thread creation

```
int pthread_create(  
    pthread_t *thread,  
    const pthread_attr_t *attr,  
    void *(*routine)(void*),  
    void *arg)
```

- `thread` thread identifier
- `routine` function that will be executed by the thread
- `arg` pointer to the argument value with which the thread function `routine()` will be executed
- `attr` use `NULL` for the time being

Reference: thread termination

A thread terminates when:

1. Thread reaches the end of its thread function, i.e., returns.
2. Thread calls

```
void pthread_exit(void *valuep)
```

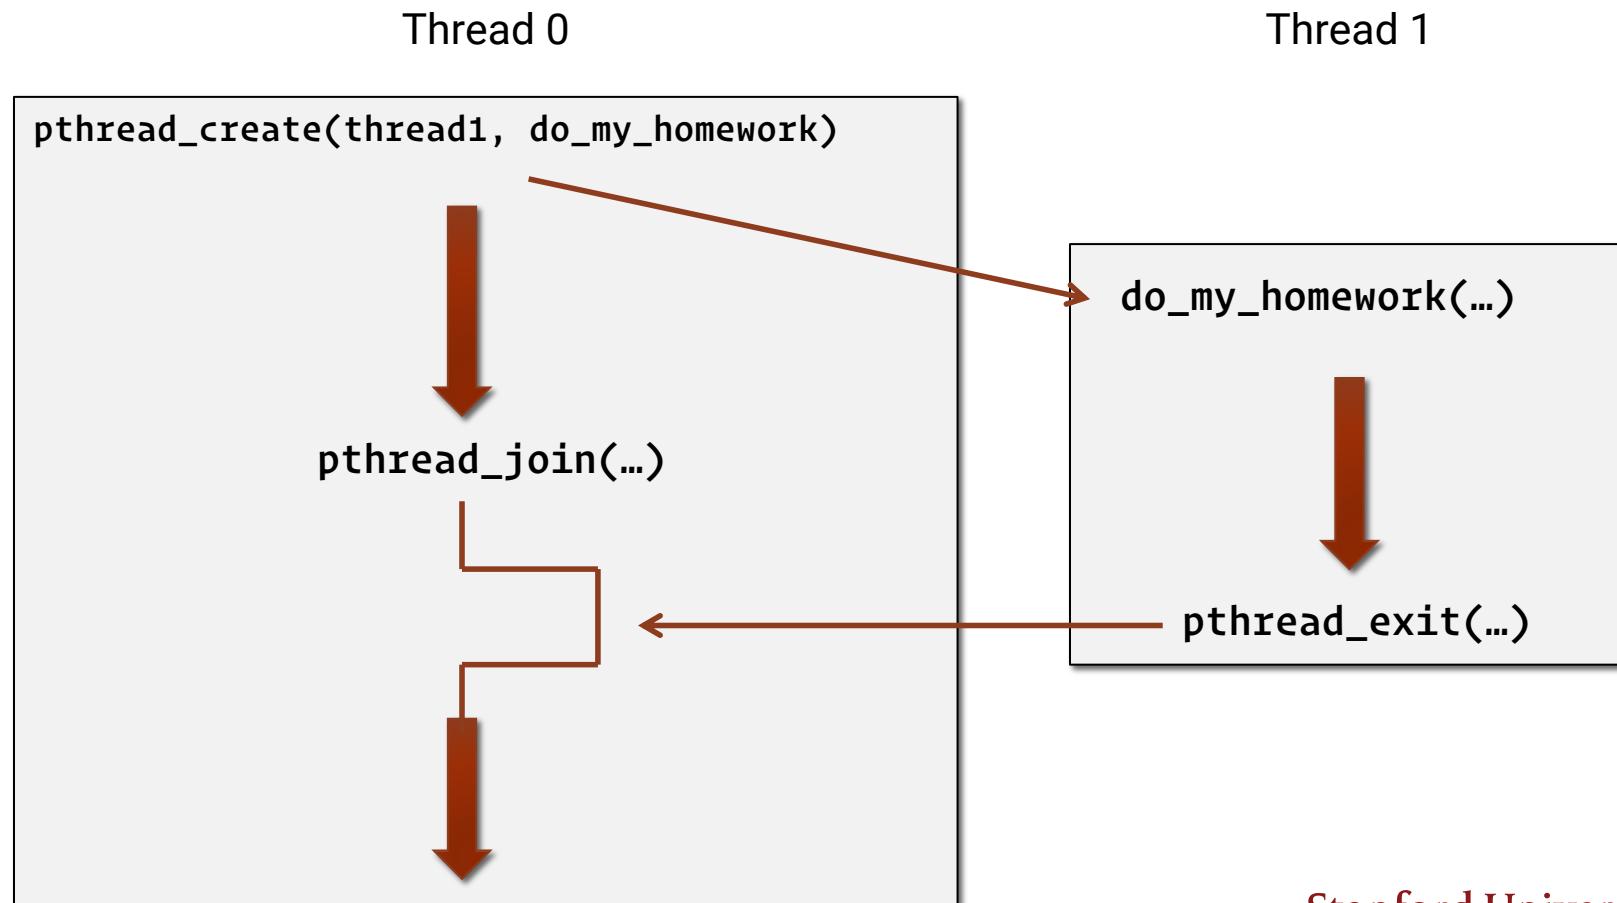
Note:

- Upon termination, a thread releases its runtime stack.
- Therefore, the pointer should point to: 1) a global variable, or 2) a dynamically allocated variable.

```
void *PrintHello(void *threadid)
{
    long tid = (long)threadid;
    long result;
    result = ComplexCalculation(tid); /* Simulates some useful calculation */
    printf("Hello World! It's me, thread #%ld. My value is %ld!\n", tid,
           result);
    pthread_exit((void *)result /*NULL is common*/);
}
```

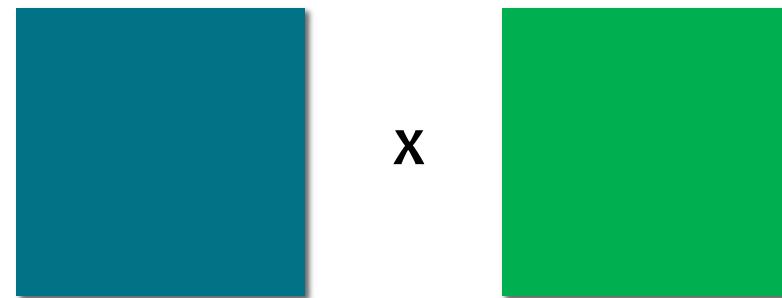
```
int pthread_join(pthread_t thread, void **valuep)
```

- Calling thread waits for thread to terminate.
- `pthread_join` is used to synchronize threads.
- `valuep` memory address where the return value of thread will be stored.



Example: matrix matrix product

- $C = A * B$
where A and B are square matrices.
- `matrix_prod.cpp` `matrix_prod_C.c`



Code snippet

```
/* Create all threads and do the computation */
for (int i = 0; i < n_thread; ++i)
    threads.push_back(thread(MatrixMult, size,
                             i * nrow_per_thread, /* row_start */
                             min(size, (i + 1) * nrow_per_thread), /* row_end */
                             ref(mat_c)));
```

```
void MatrixMult(int size, int row_start, int row_end, vector<float> &mat_c)
{
    printf("Processing row %3d to %3d\n", row_start, row_end - 1);

    for (int i = row_start; i < row_end; ++i)
        for (int j = 0; j < size; ++j)
        {
            float c_ij = 0.;

            for (int k = 0; k < size; ++k)
            {
                c_ij += MatA(i, k) * MatB(k, j);
            }

            mat_c[i * size + j] = c_ij;
        }
}
```

Thread coordination



The risks of multi-threaded programming

- Let us assume that a well-known bank company has asked you to implement a multi-threaded code to perform bank transactions.
- You start with the modest goal of allowing deposits.
- Clients deposit money and the amount gets credited to their accounts.
- As a result of having multiple threads running concurrently the following can happen:

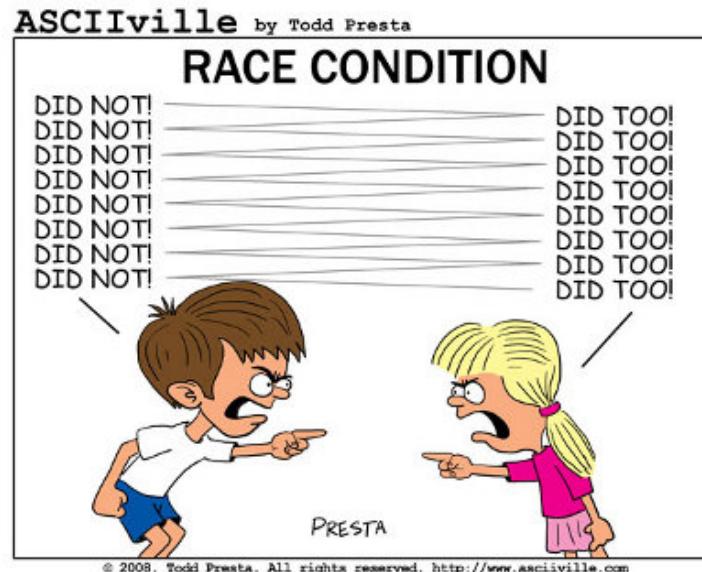


A parallel bank deposit

Thread 0	Thread 1	Balance
Client requests a deposit	Client requests a deposit	\$1000
Check current balance = \$1000	Check current balance = \$1000	
Ask for deposit amount = \$100	Ask for deposit amount = \$300	
	Compute new balance = \$1300	
Compute new balance = \$1100	Write new balance to account	\$1300
Write new balance to account		\$1100

Race condition

- Although the correct balance should be \$1,400, it is \$1,100.
- The problem is that many operations “take time” and can be interrupted by other threads attempting to modify the same data.
- This is called a race condition: the final result depends on the precise order in which the instructions are executed.



Race condition

- This issue is addressed using mutexes (mutex): mutual exclusion.
- They ensure that certain common pieces of data are accessed and modified by a single thread.
- This problem typically occurs when you have a sequence like:

READ/WRITE, or
WRITE/READ

performed by different threads.



Thread 0 wants to
add new to-do item.



Thread 0 closes lock.
No other thread can open
the lock.



Thread 0 is done with the
to-do list.
It opens the lock.



Thread 1 wants to open
the lock.
It has to wait.

Thread 1 can close the lock
and access the to-do list.

Mutex

- A mutex can only be in two states: locked or unlocked.
- Once a thread locks a mutex:
 - Other threads attempting to lock the same mutex are blocked
 - Only the thread that initially locked the mutex has the ability to unlock it.
- This allows to protect regions of code: only one thread at a time can execute that code.



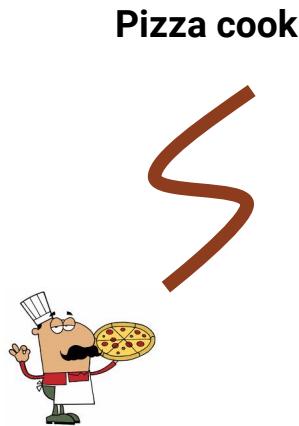
Typical usage

Mutex use:

- Create and initialize a mutex variable.
- Threads attempt to lock the mutex.
- Only one succeeds, and that thread owns the mutex.
- The owner thread performs some set of actions.
- The owner unlocks the mutex.
- Another thread acquires the mutex and repeats the process.



Pizza restaurant



Pizza delivery team



- Checks the addresses for customers
- Deliver pizza

Go back;
check if there are
orders left.

Pizza restaurant code

mutex_demo.c

```

while (1)
{
    pthread_mutex_lock(&mutex);

    if (global_task_list != NULL)
    {
        my_task = global_task_list;                      /* Pop a task */
        Wait();   /* This causes several threads to use the same task pointer */
        global_task_list = global_task_list->next; /* Move forward by one task */
        pthread_mutex_unlock(&mutex);

        /* At this point, the thread delivers the pizza. */
        /* The mutex is unlocked so that other threads can work. */
        printf("Thread %ld: %s\n", thread_id, my_task->work_to_do);
        Delivery();

        /* Work is done */
        free(my_task->work_to_do); /* Free task data */
        free(my_task);             /* Free task itself */
    }
    else
    {
        pthread_mutex_unlock(&mutex); /* Don't forget to unlock the mutex */
        pthread_exit(NULL);
    }
}

```

Reference: summary of key functions

```
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
```

- Initialization of mutex; choose NULL for attr.

```
int pthread_mutex_init(pthread_mutex_t *mutex,  
                      const pthread_mutexattr_t *attr)
```

- Destruction of mutex.

```
int pthread_mutex_destroy(pthread_mutex_t *mutex)
```

- Locks a mutex; blocks if another thread has locked this mutex and owns it.

```
int pthread_mutex_lock(pthread_mutex_t *mutex)
```

- Unlocks mutex; after unlocking, other threads get a chance to lock the mutex.

```
int pthread_mutex_unlock(pthread_mutex_t *mutex)
```

Deadlock

Another strange case (but common!) of parallel computing.

Let's assume:

The code
will never
reach this
point

Thread 0 Locks *mutex0*
Thread 1 Locks *mutex1*
Thread 0 Does some work
Thread 1 Does some work
Thread 0 Locks *mutex1*
Thread 1 Locks *mutex0*
→ Thread 0 Work requiring lock on both mutexes
Thread 1 Work requiring lock on both mutexes
Thread 0 Unlocks all mutexes
Thread 1 Unlocks all mutexes



Solution: threads always lock mutexes in the same order.

Condition variables

Are we there yet?

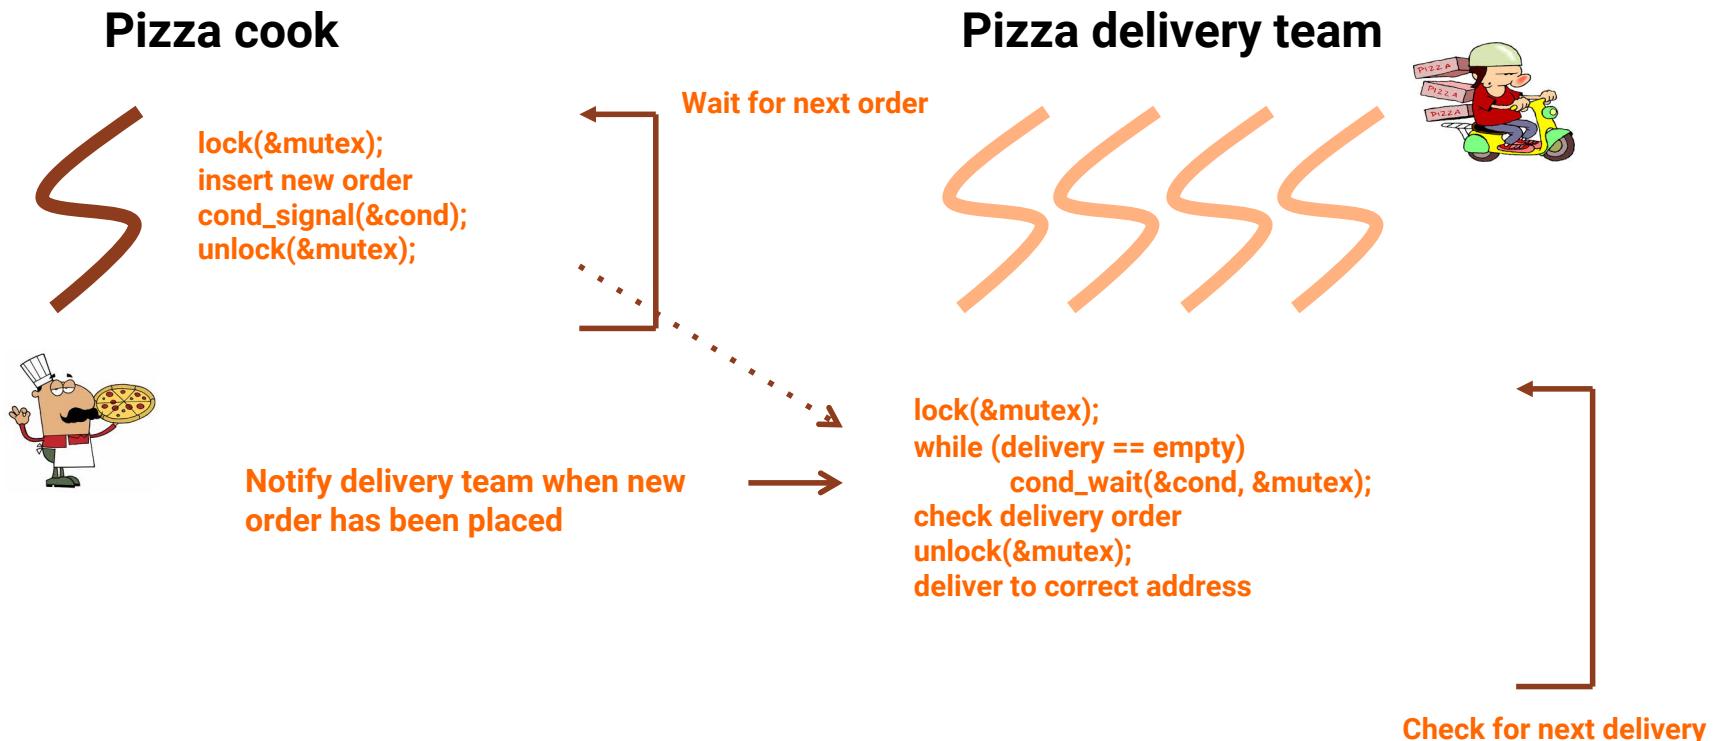


- There is a common situation in computing: you have a pool of threads that are ready to perform tasks, but there is nothing to do yet. So they have to wait until some work becomes available.
- Instead of constantly checking the to-do list, they can simply wait and be awoken when a certain condition is met.



Waiting on a condition

- Application: pizza delivery boys waiting on pizza cook.
- This is the producer-consumer model.



Condition variable example

`cond_var.c`

```
while (tasks_to_do--)
{
    RandomWait();
    pthread_mutex_lock(&mutex);
    while (global_task_list == NULL)
    {
        printf("Consumer thread %ld: waiting for next task\n", thread_id);
        pthread_cond_wait(&cond, &mutex);
        /* Mutex unlocks while waiting.
         * Locks when returning from pthread_cond_wait(). */
    }

    my_task = global_task_list; /* Task that thread is going to process */
    global_task_list = global_task_list->next; /* Next task */
    pthread_mutex_unlock(&mutex);

    printf("Consumer thread %ld: grabbed task %s\n", thread_id,
           my_task->work_to_do);
    printf("Consumer thread %ld: mutex unlocked\n", thread_id);

    /* At this point, the thread delivers the pizza. */
    /* The mutex is unlocked so that other threads can work. */
    free(my_task->work_to_do);
    free(my_task);
}

pthread_cond_signal(&cond); /* Wake up a consumer. */
```

Condition variable

- Requires a condition variable and a mutex.
- The mutex is used to protect the access to the condition variable.

```
int pthread_cond_wait(pthread_cond_t *cond,  
                      pthread_mutex_t *mutex)
```

Recommended
scenario

```
pthread_mutex_lock(&mutex);  
while (!condition_ready())  
    pthread_cond_wait(&cond, &mutex);  
access_modify_shared_data();  
pthread_mutex_unlock(&mutex);
```

When `pthread_cond_wait` is called:

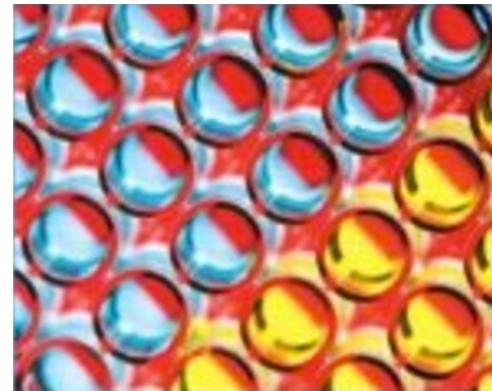
- The thread releases the mutex `mutex`.
- The thread waits until a signal is sent.
- Upon receiving a signal, the thread locks `mutex` and proceeds.

Condition signal

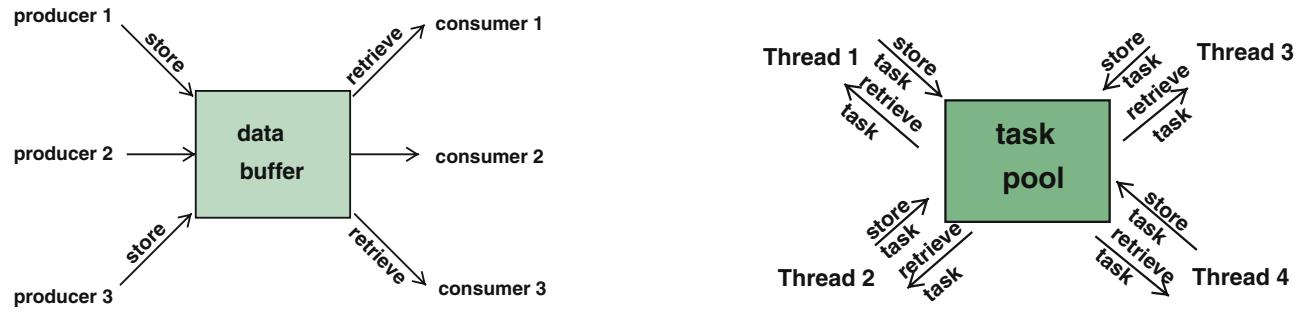
```
int pthread_cond_signal(pthread_cond_t *cond)
```

- Wakes up a single thread waiting on the variable cond.
- Usually, around `pthread_cond_signal`, there is code that reads/modifies the condition and needs to be protected using a mutex.
- `pthread_cond_signal` is commonly placed inside the mutex-protected block, although this is not necessary.

Parallel patterns and other features of Pthreads

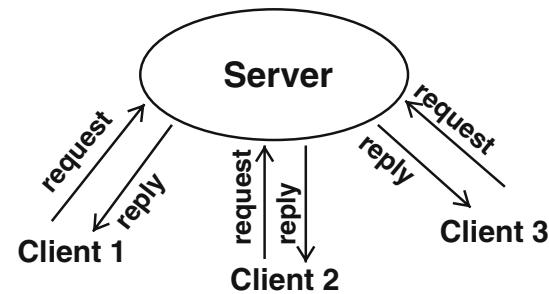


Example of useful parallel programming patterns



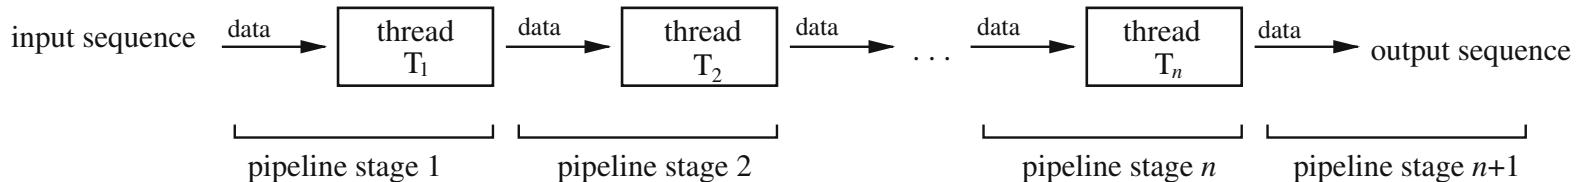
Producer consumer

Task pool

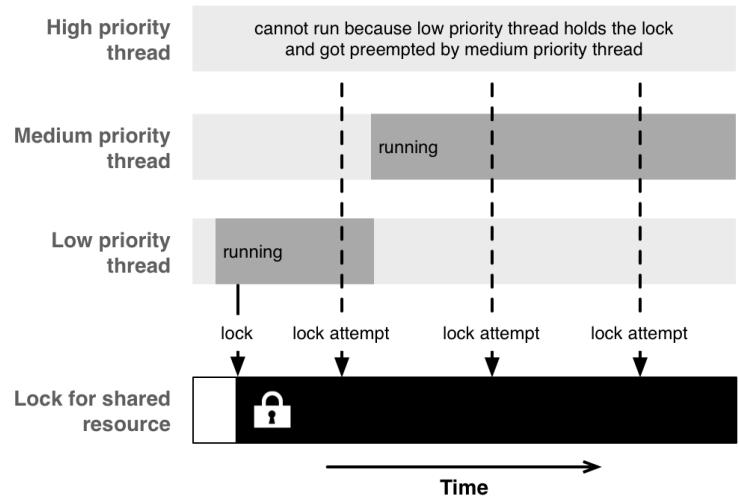


Client server

Pipeline



Other features



- Thread scheduling
 - › Implementations will differ on how threads are scheduled to run. In most cases, the default mechanism is adequate.
 - › The Pthreads API provides routines to explicitly set thread scheduling policies and priorities which may override the default mechanisms.
- Thread-specific data: keys
 - › To preserve stack data between calls, you can pass it as an argument from one routine to the next, or else store the data in a global variable associated with a thread.
 - › Pthreads provides another, possibly more convenient and versatile, way of accomplishing this, through keys.
- Priority inversion problems; priority ceiling and priority inheritance
- Thread cancellation
- Barriers; not always available
- Thread safety