

CME 213

SPRING 2019

Eric Darve

Shared Memory Sorting

Sorting

- Sorting is a key algorithm in many engineering problems.
- They are both ubiquitous and difficult to implement.
- Sorting is significantly more difficult in parallel than in sequential.
- Algorithms well-suited for sequential sorts are often ill-suited on a parallel machine.
- A fascinating topic!

Quicksort and mergesort

Two great sorting algorithms:

- Occupy a prominent place in world's computational infrastructure.
- Quicksort honored as one of **top 10 algorithms of 20th** century in science and engineering.
- Quicksort

Quicksort

- Most common algorithm.
- Simple, low overhead, optimal **average** complexity.
- Divide and conquer approach.
- Divide step:
 - › Choose a pivot x .
 - › Separate sequence into 2 sub-sequences with all elements smaller than x and greater than x .
- Conquer step:
 - › Sort the two subsequences.

See Python code `sort.py`

```
def quicksort(A, l, u):  
    if l < u-1:  
        x = A[l]  
        s = l  
        for i in range(l+1, u):  
            if A[i] <= x:  
                s = s+1  
                A[s], A[i] = A[i], A[s]  
        A[s], A[l] = A[l], A[s]  
        quicksort(A, l, s)  
        quicksort(A, s+1, u)
```

Parallel Quicksort

- Once we have the two sub-lists, we can process each sublist in parallel.
- This works well once the number of sublists is greater than the number of threads.
 - › One difficulty is that the length of the sublists can be very different, leading to parallel load-imbalance.
- The creation of the sub-lists in parallel is not easy.
- We split the list into smaller chunks.
- Each chunks is rearranged (pivot rearrangement).
- Then, we compute the global rearrangement.

First Step

P_0	P_1	P_2	P_3	P_4																
7	13	18	2	17	1	14	20	6	10	15	9	3	16	19	4	11	12	5	8	
pivot=7																				pivot selection
P_0	P_1	P_2	P_3	P_4																
7	2	18	13	1	17	14	20	6	10	15	9	3	4	19	16	5	12	11	8	
																				after local rearrangement
7	2	1	6	3	4	5	18	13	17	14	20	10	15	9	19	16	12	11	8	
																				after global rearrangement

pivot selection

after local rearrangement

after global rearrangement

Second Step

P_0	P_1	P_2	P_3	P_4															
7	2	1	6	3	4	5	18	13	17	14	20	10	15	9	19	16	12	11	8
pivot=5							pivot=17												
pivot selection																			
P_0	P_1	P_2	P_3	P_4															
1	2	7	6	3	4	5	14	13	17	18	20	10	15	9	19	16	12	11	8
after local rearrangement																			
1	2	3	4	5	7	6	14	13	17	10	15	9	16	12	11	8	18	20	19
after global rearrangement																			

pivot selection

after local rearrangement

after global rearrangement

Third Step

P_0	P_1	P_2	P_3	P_4															
1	2	3	4	5	7	6	14	13	17	10	15	9	16	12	11	8	18	20	19
pivot=11																	pivot selection		
P_0	P_1	P_2	P_3	P_4															
1	2	3	4	5	6	7	10	13	17	14	15	9	8	12	11	16	18	19	20
							10	9	8	12	11	13	17	14	15	16			
																	after local rearrangement		
							10	9	8	12	11	13	17	14	15	16			
																	after global rearrangement		

pivot selection

after local rearrangement

after global rearrangement

Fourth Step

																				</									

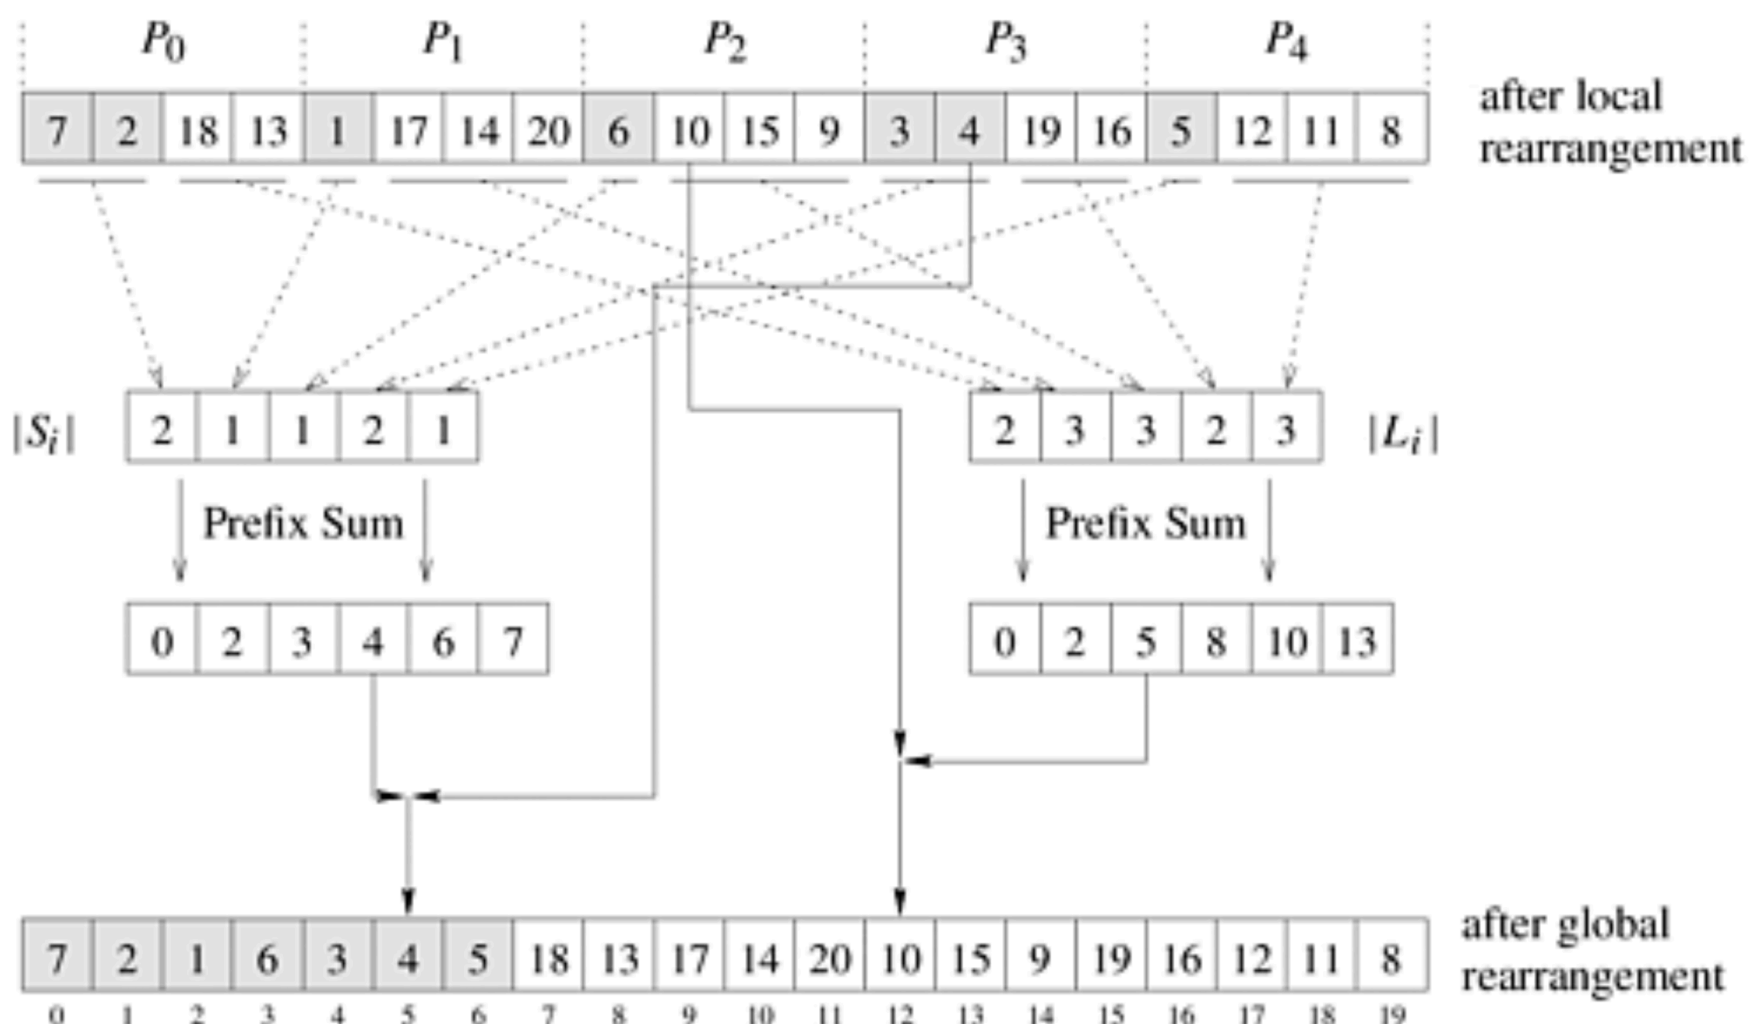

after local rearrangement

P_0					P_1		P_2					P_3					P_4			Solution
1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	

Solution

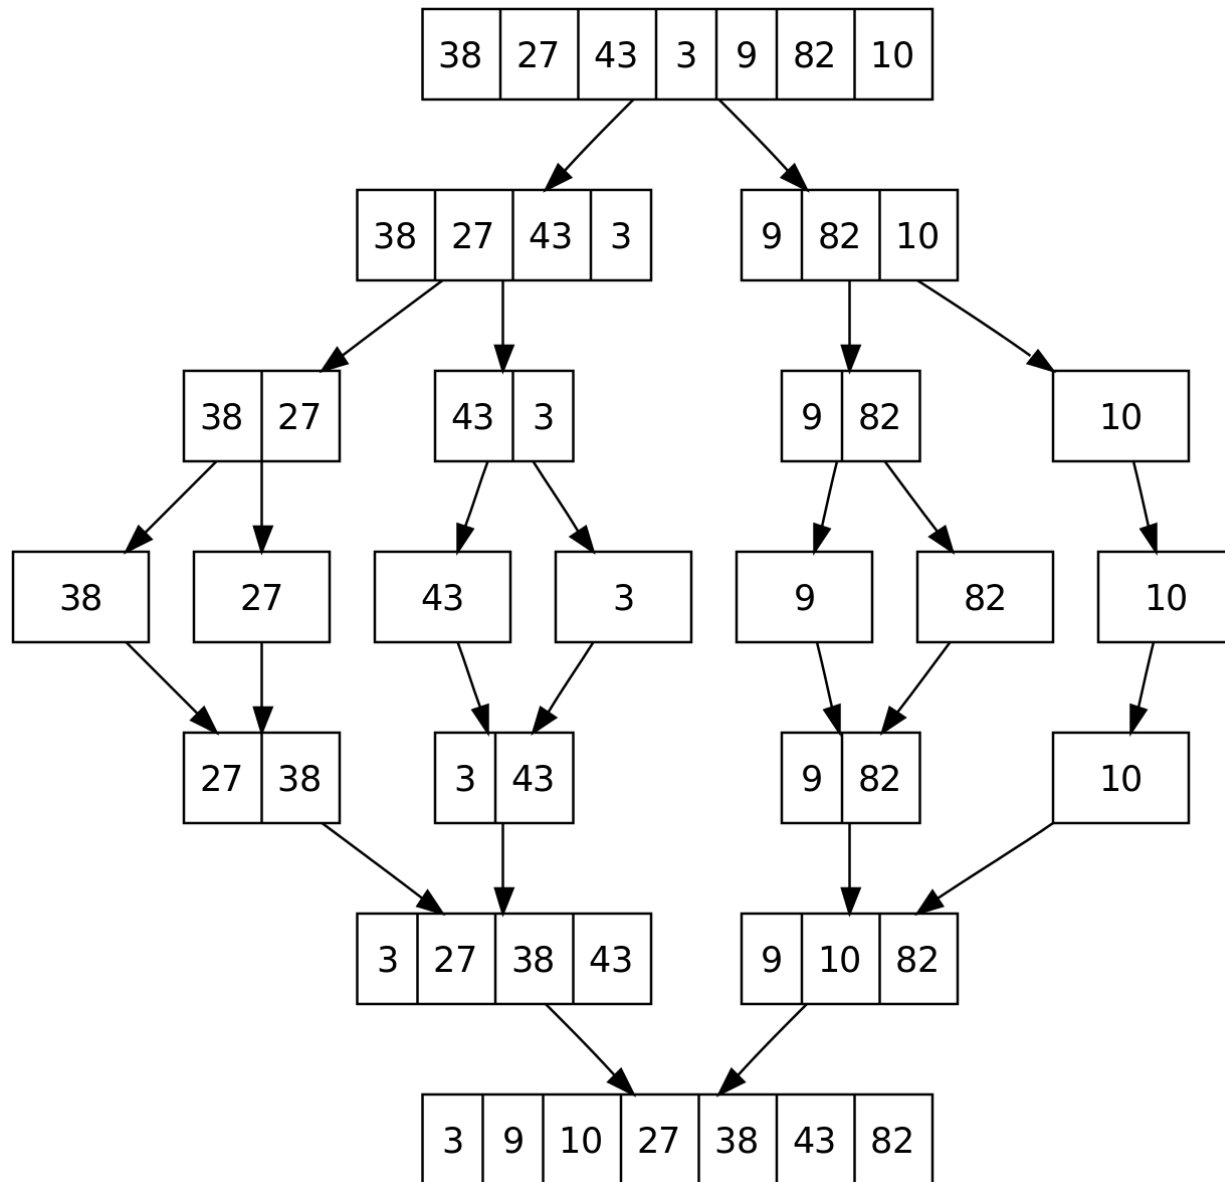
Prefix sum used to computed the starting addresses for the global rearrangement.

Load imbalance because of sublist sizes and number of processes assigned to each sublist



Mergesort


- Another popular and classic algorithm.
- We first subdivide the list into n sub-lists (each with one element).
- Then sub-lists are progressively merged to produce larger ordered sub-lists.
- Merge sort



Parallel implementation

- When there are many sub-lists to merge, the parallel implementation is straightforward: assign each sub-list to a thread.
- When we get few but large sub-lists, the parallel merge becomes more difficult.
- In that case, we need a way to subdivide the merge into several smaller merges that can be done concurrently.


Divide and conquer parallel merge




Big merge replaced by 2 smaller parallel merges

Bucket and sample sort


- Bucket sort is a simpler parallel algorithm.
- Assume we have a sequence of integers in the interval $[a,b]$.
- Split $[a,b]$ into p sub-intervals.
- Move each element to the appropriate bucket (prefix sum required again).
- Sort each bucket in parallel.
- Simple and efficient!
- Radix sort
- Problem: how should we split the interval? This process may lead to intervals that are unevenly filled.
- Improved version: sample (or splitter) sort.




Initial element distribution



Local sort & sample selection




Sample combining



Global splitter selection

Prefix sum required here to partition the input sequence



Final element assignment

Final step is similar to bucket/radix sort

Performance

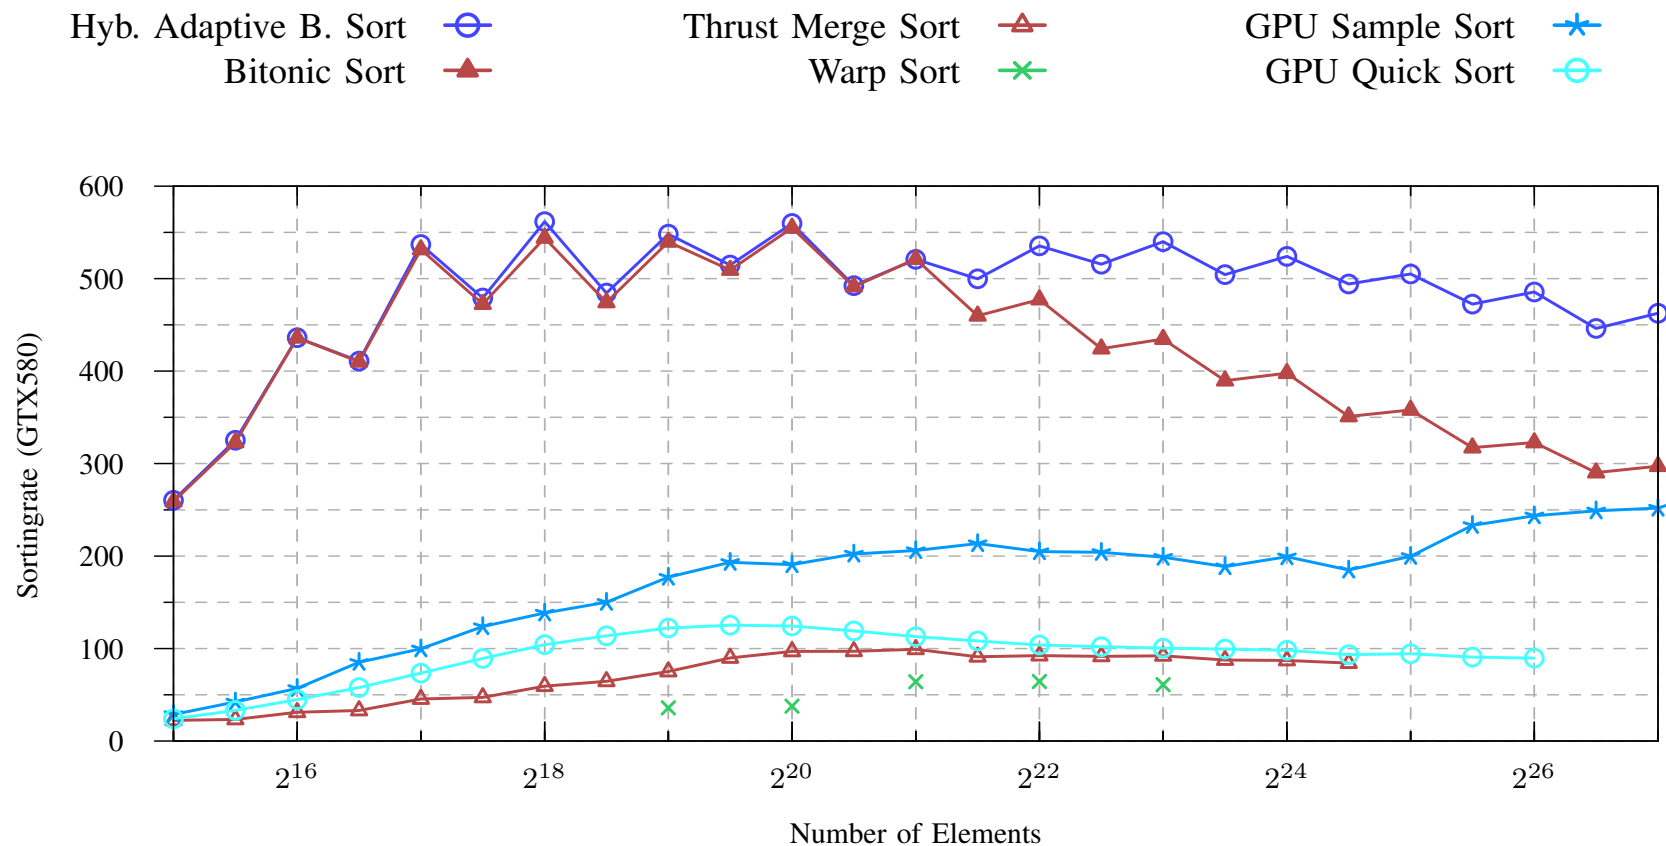
- Sample sort is one of the fastest parallel sorting algorithm.
- One potential bottleneck is sorting the samples (global splitter selection), if this is done sequentially.

Sorting networks

- This is a very special class of sorting algorithms.
- In sorting networks, the sequence of operations (called compare-and-exchange COEX) is independent of the data!
- That is, we perform a deterministic sequence of COEX operations that result in a sorted sequence!
- These algorithms are more complicated but can result in better performance in some cases.
- One of their advantages is that they are very regular compared to the other sorting algorithms.


A novel sorting algorithm for many-core architectures based on adaptive bitonic sort

H. Peters, O. Schulz-Hildebrandt, N. Luttenberger



Bitonic sequence

- What is a bitonic sequence?
- A sequence of n numbers such that:
 - › The first part is ↗ and the second part is ↘
 - › The first part is ↘ and the second part is ↗




Bitonic merge step

- Consider a bitonic sequence E of length n .
- Define:

$$L(E) = \left(\min(E_0, E_{n/2}), \min(E_1, E_{n/2+1}), \dots, \min(E_{n/2-1}, E_{n-1}) \right)$$


$$U(E) = \left(\max(E_0, E_{n/2}), \max(E_1, E_{n/2+1}), \dots, \max(E_{n/2-1}, E_{n-1}) \right)$$

- These two subsequences are bitonic.
- Any element a in $L(E)$ is smaller than any element b in $U(E)$.



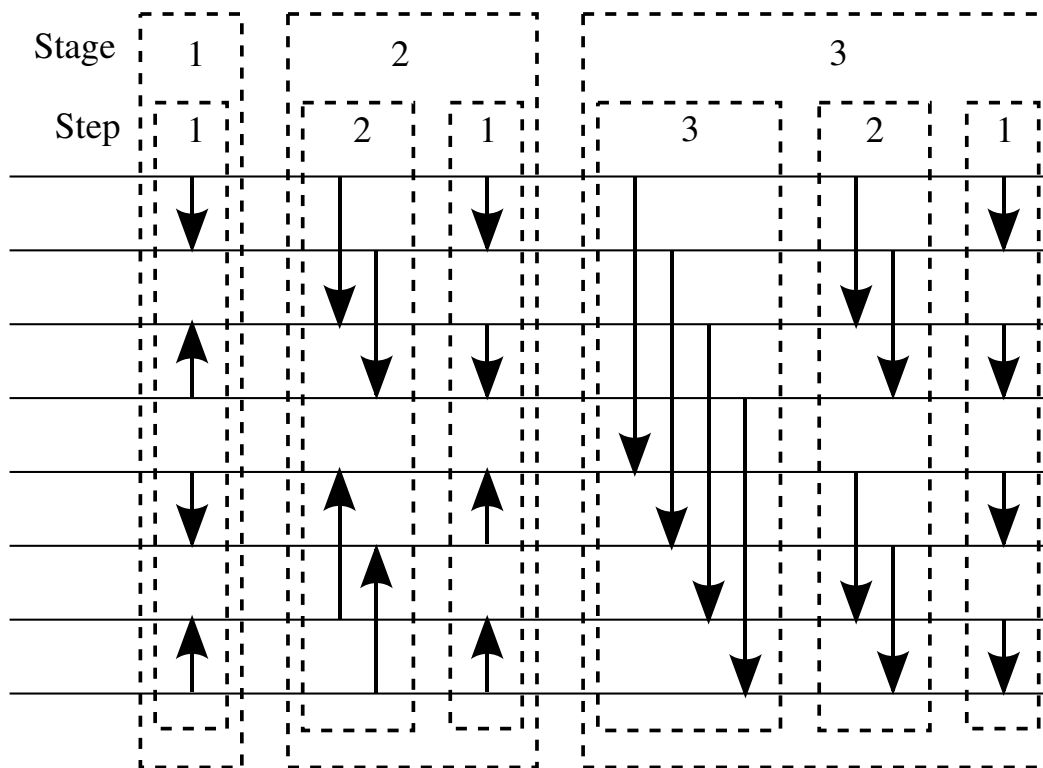
Bitonic merging network

- Start from a bitonic sequence.
- Repeat previous merging process to get a sorted sequence.



Bitonic sorting network

- Start from a sequences of length **2** and sort them.
- Then, sort sequences of lengths **4, 8, 16**, etc, until the whole sequence is sorted.



See Python code `bitonic_sort.py`

Bitonic

Computational cost

- This method has cost $O(n \ln^2 n)$.
- This cost can be reduced using an **adaptive bitonic sort**.
- In this variant, the bitonic **merge** consists of:
 1. Compute how intervals must be rearranged in $\log(n)$ steps (for a sequence of size n).
 2. Rearrange the intervals to create two bitonic sub-sequences in $\log(n)$ steps using a bitonic tree; $\log(n)$ pointers need to be swapped instead of copying data.
- With this algorithm, each merge costs only $O(n)$.
- Total cost is $O(n \ln n)$.
- Other optimizations allow making this algorithm much faster than a sample sort on GPUs.

Bitonic sort lab

Different codes are provided:

- **bitonic_sort_seq.cpp**
Reference sequential implementation
- **bitonic_sort_lab.cpp**
Open this code to start the exercise
- **bitonic_sort.cpp**
Solution with OpenMP

Parallel steps

There are two main parallelization strategies:

1. When we have many bitonic sequences, we can parallelize over the j loop
2. When the bitonic sequences are long, we can parallelize the bitonic merge step (min/max loop).