

CME 213

SPRING 2019

Eric Darve

Review

“Secret” behind GPU performance: simple cores but a large number of them; even more threads can exist live on the hardware (10k–20k threads live).

Important performance bottleneck: data traffic between memory and register files

Programming on a GPU:

- Warp = 32 threads
- Block = group of threads running on one SM
- SM runs an integer number of thread blocks
- Grid = set of blocks representing the entire data set

Code example

- Let's start with a very basic example
`firstProgram.cu`
- Log onto GCP!

```
$ gcloud compute ssh hw3  
$ gcloud compute scp * hw3:~/.
```

Compile with:

```
$ make
```

or:

```
$ nvcc --compiler-options -Wall -arch=sm_35 -o test test.cu
```

Code overview

```
int* d_output;
cudaMalloc(&d_output, sizeof(int) * N);
kernel<<<1, N>>>(d_output);
vector<int> h_output(N);
cudaMemcpy(&h_output[0], d_output, sizeof(int) * N, cudaMemcpyDeviceToHost);
for(int i = 0; i < N; ++i)
    printf("Entry %3d, written by thread %2d\n", h_output[i], i);
cudaFree(d_output);
```

Function to execute

Function arguments



```
kernel<<<1, N>>>(d_output);
```

Number of threads to use

Device kernels

```
__device__ __host__  
int f(int i) {  
    return i*i;  
}
```

```
__global__  
void kernel(int* out) {  
    out[threadIdx.x] = f(threadIdx.x);  
}
```

global host device

What are these mysterious keywords?

`__global__` kernel will be

- Executed on the device
- Callable from the host

`__host__` kernel will be

- Executed on the host
- Callable from the host



Compiled for host

`__device__` kernel will be

- Executed on the device
- Callable from the device only



Compiled for device

Run the following!

```
./firstProgram 1  
./firstProgram 32  
./firstProgram 1024  
./firstProgram 1025
```

Try with arguments:

1, 2, 3, 4, 16, 32, 64, 1024, **1025**


```
darve@hw3:~$ ./getProperties
CUDA Device Query...
There are 1 CUDA devices.

CUDA Device #0
Major revision number:      3
Minor revision number:      7
Name:                       Tesla K80
Total global memory:        11996954624
Total shared memory per block: 49152
Total registers per block:  65536
Warp size:                  32
Maximum memory pitch:       2147483647
Maximum threads per block:  1024
Maximum dimension 0 of block: 1024
Maximum dimension 1 of block: 1024
Maximum dimension 2 of block: 64
Maximum dimension 0 of grid: 2147483647
Maximum dimension 1 of grid: 65535
Maximum dimension 2 of grid: 65535
Clock rate:                  823500
Total constant memory:       65536
Texture alignment:           512
Concurrent copy and execution: Yes
Number of multiprocessors:   13
Kernel execution timeout:    Yes
ECC enabled:                 Yes
```

Thread hierarchy

Two levels of hierarchy:

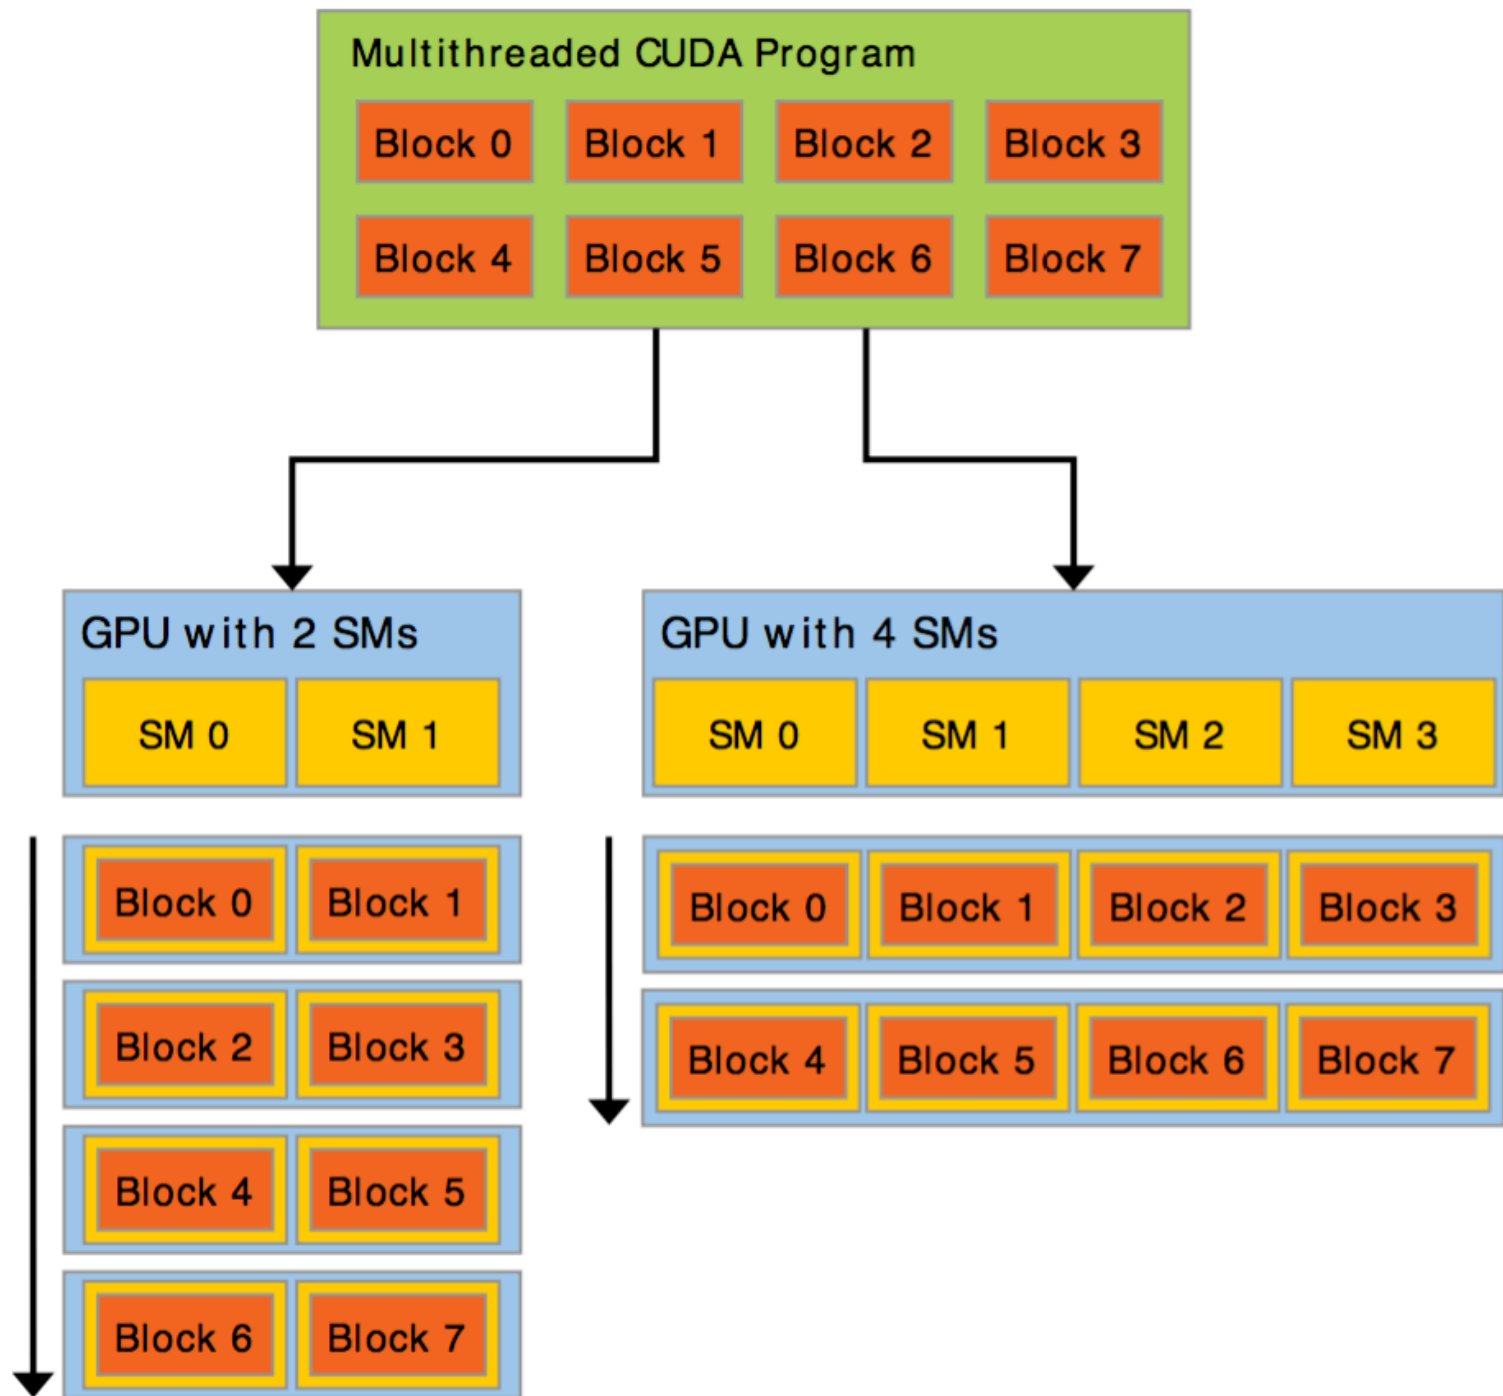
- **Thread block:** a block is loaded entirely on an SM. This means that a block of threads cannot require more resources than available, for example to store register variables.
- **Grid of blocks:** for large problems, we can use multiple blocks. A very large number of blocks can be allocated so that we can execute a kernel over a very large data set.

Problem decomposition

Decompose your problem into:

- Small blocks so that data can fit in SM.
- Typically the number of threads per block is around 256–512.
- Create a grid of blocks so that you can process the entire data.

```
Total shared memory per block: 49152
Total registers per block:      65536
Warp size:                      32
Maximum memory pitch:          2147483647
Maximum threads per block:     1024
Maximum dimension 0 of block:  1024
Maximum dimension 1 of block:  1024
Maximum dimension 2 of block:   64
Maximum dimension 0 of grid:    2147483647
Maximum dimension 1 of grid:    65535
Maximum dimension 2 of grid:    65535
```



Block execution

- The device will start by loading data on each SM to execute the blocks.
- At least one block must fit on an SM, e.g., there should be enough memory for register variables.
- The SM will load up as many blocks as possible until it runs out of resources.
- Then it will start executing the blocks, that is all threads will execute the kernel.
- Once all threads in a block are done with the kernel, another block gets loaded in the SM.
- And so on until the device runs out of blocks to execute.

Dimensions

Blocks and grids can be 1D, 2D or 3D.

Their dimensions is declared using:

```
dim3 threadsPerBlock(Nx);
```

```
dim3 numBlocks(Mx);
```

```
dim3 threadsPerBlock(Nx, Ny);
```

```
dim3 numBlocks(Mx, My);
```

```
dim3 threadsPerBlock(Nx, Ny, Nz);
```

```
dim3 numBlocks(Mx, My, Mz);
```

Thread index example:

```
int col = blockIdx.x * blockDim.x + threadIdx.x;
```

```
int row = blockIdx.y * blockDim.y + threadIdx.y;
```

Block size



**Thread ID in
block**



Block ID



addMatrices.cu

```
__global__  
void Add(int n, int* a, int* b, int* c) {  
    int i = blockIdx.x * blockDim.x + threadIdx.x;  
    int j = blockIdx.y * blockDim.y + threadIdx.y;  
    if(i < n && j < n) {  
        c[n*i + j] = a[n*i + j] + b[n*i + j];  
    }  
}
```

```
dim3 threads_per_block(2, n_thread);  
int blocks_per_grid_x = (n + 2 - 1) / 2;  
int blocks_per_grid_y = (n + n_thread - 1) / n_thread;  
dim3 num_blocks(blocks_per_grid_x, blocks_per_grid_y);  
Initialize<<<num_blocks, threads_per_block>>>(n, d_a, d_b);  
Add<<<num_blocks, threads_per_block>>>(n, d_a, d_b, d_c);
```

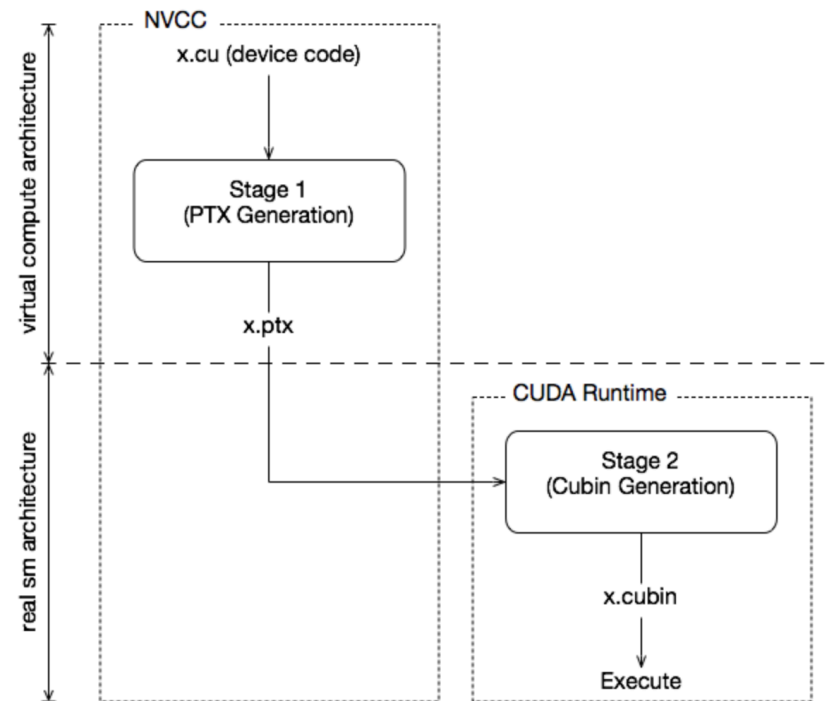
Compiler options and other tricks

- g Debug on the host
- G Debug on the device (CUDA-gdb, Nsight Eclipse Edition)
- pg Profiling info for use with gprof (Linux only)
- Xcompiler Options for underlying gcc compiler
- O Optimization level

More on profiling: **nvprof**

<https://docs.nvidia.com/cuda/cuda-compiler-driver-nvcc/index.html>

GPU architecture



nvcc produces different types of codes:

- **Architecture:** a virtual architecture with certain capabilities. Only declare the capabilities your code is using (lowest possible number)
- **Real target GPU:** compilation of the code for a specific GPU (many can be specified)

The architecture and GPU code types can be specified:

```
$ nvcc code.cu --gpu-architecture=compute_35 --gpu-code=sm_35
```

```
$ nvcc code.cu -arch=sm_35
```

Profiling


- The Visual Profiler (nvvp) is the simplest way to profile your code. You get a graphical interface that helps navigate the information.
- You can also use a command line tool (good for clusters): nvprof.
- This can also be used to collect information on a cluster to be visualized using nvvp on your local computer.

<https://docs.nvidia.com/cuda/profiler-users-guide/index.html>

Run

```
$ nvprof ./addMatrices -n 4000
```

Visual Profiler Timeline



Nvprof output

==2380== Profiling application: ./addMatrices -n 4000

==2380== Profiling result:

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	80.22%	11.119ms	1	11.119ms	11.119ms	11.119ms	[CUDA memcpy DtoH]
	12.76%	1.7687ms	1	1.7687ms	1.7687ms	1.7687ms	Add(int, int*, int*, int*)
	7.02%	972.60us	1	972.60us	972.60us	972.60us	Initialize(int, int*, int*)
API calls:	97.51%	587.99ms	3	196.00ms	183.78us	587.61ms	cudaMalloc
	1.89%	11.374ms	1	11.374ms	11.374ms	11.374ms	cudaMemcpy
	0.46%	2.7752ms	2	1.3876ms	982.79us	1.7924ms	cudaDeviceSynchronize
	0.07%	450.42us	1	450.42us	450.42us	450.42us	cuDeviceTotalMem
	0.05%	287.33us	96	2.9920us	129ns	132.76us	cuDeviceGetAttribute
	0.01%	86.226us	2	43.113us	19.461us	66.765us	cudaLaunchKernel
	0.00%	18.389us	1	18.389us	18.389us	18.389us	cuDeviceGetName
	0.00%	2.7500us	1	2.7500us	2.7500us	2.7500us	cuDeviceGetPCIBusId
	0.00%	2.5890us	3	863ns	198ns	1.6910us	cuDeviceGetCount
	0.00%	1.8980us	2	949ns	239ns	1.6590us	cuDeviceGet
	0.00%	810ns	2	405ns	269ns	541ns	cudaGetLastError
	0.00%	285ns	1	285ns	285ns	285ns	cuDeviceGetUuid

cuda-memcheck

Several tools available to verify your code.

<https://docs.nvidia.com/cuda/cuda-memcheck/index.html>

4 tools accessible with `cuda-memcheck` command:

- Memcheck Memory access checking
- Initcheck Global memory initialization checking
- Racecheck Shared memory hazard checking
- Synccheck Synchronization checking

Usage examples:

```
$ cuda-memcheck ./memcheck_demo
```

```
$ cuda-memcheck --leak-check full ./memcheck_demo
```

GPU Optimization!



Top factors to consider

- Memory: data movement!
- Occupancy and concurrency:
“hide latency through concurrency”
- Control flow: branching

Memory

- The number one bottleneck in scientific applications are not flops!
- Flops are plentiful and abundant.
- The problem is data starvation.
- Computing units are idle because they are waiting on data coming from / going to memory.

Memory access

- Similar to a regular processor but with some important differences.
- Caches are used to optimize memory accesses: L1 and L2 caches.
- Cache behavior is complicated and depends on the compute capability of the card (that is the type of GPU you are using).
- Generations (compute capability): 2.x, 3.x, 5.x, and 6.x (no 4.x).
- For this class, we have access to **3.7** hardware.
- We will focus on this for our discussion.

Cache

- There are two main levels of cache: L1 and L2.
- L2 is shared across SMs.
- Each L1 is assigned to only one SM.
- When reading data from global memory, the device can only read entire cache lines.
- There is no such thing as reading a single float from memory.
- To read a single float, you read the entire cache line and “throw out” the data you don’t need.

Cache lines

- L1: 128-byte lines = 32 floats
- L2: 32-byte segments = 8 floats

How is data read from memory then?

- Remember how the hardware works.
- All threads in a warp execute the same instruction at the same time.
- Take 32 threads in a warp.
- Each thread is going to request a memory location.
- Hardware groups these requests into a number of cache line requests.
- Then data is read from memory.
- At the end of the day, what matters is not how much data threads are requesting from memory, it's how many memory transactions (cache lines) are performed!

Example


```
__global__ void offsetCopy(float* odata, float* idata, int offset) {  
    int xid = blockIdx.x * blockDim.x + threadIdx.x + offset;  
    odata[xid] = idata[xid];  
}
```

offset = 0

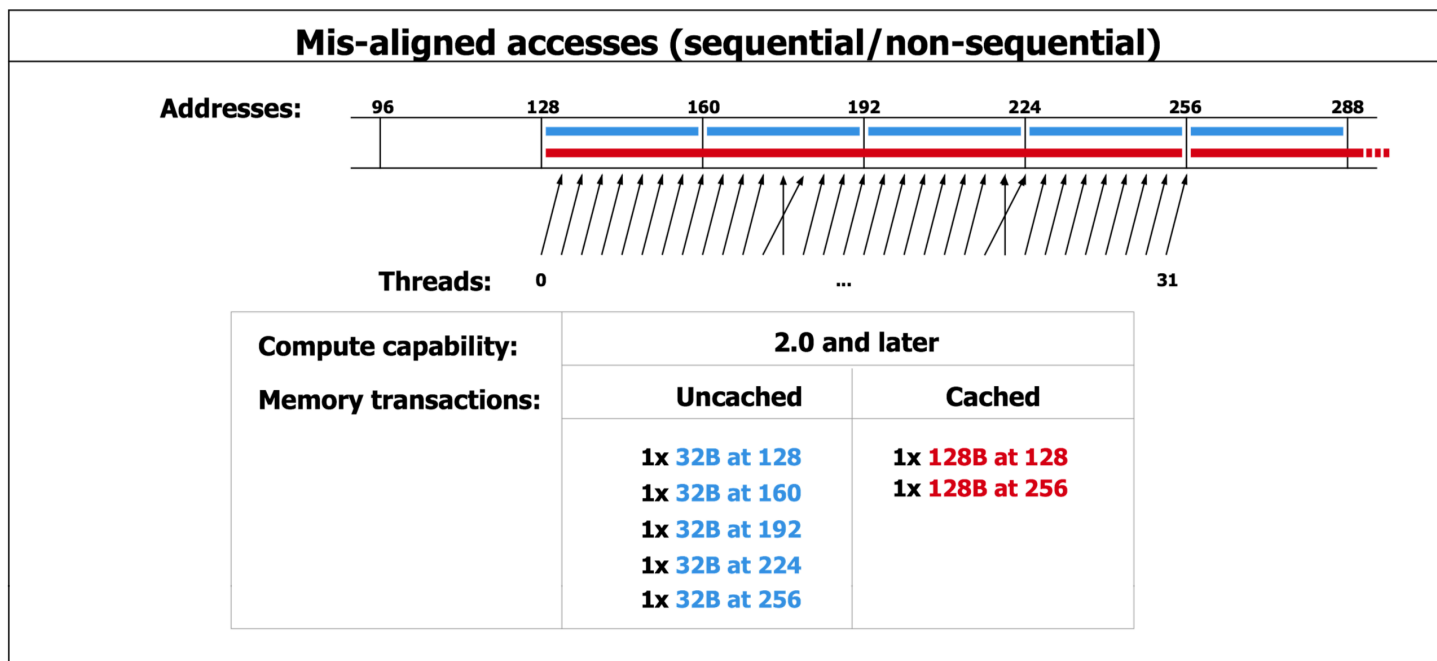
This is the best case.

All memory accesses are coalesced into a single memory transaction.

One memory transaction




Misaligned by 1



offset = 1

Strided access

```
__global__ void stridedCopy(float* odata, float* idata, int stride) {  
    int xid = blockIdx.x * blockDim.x + threadIdx.x + offset;  
    odata[xid] = idata[stride*xid];  
}
```



Stride is 2.

Effective bandwidth of application reduced by 2.

Efficiency of reads

Perfectly coalesced access: one memory transaction

Bandwidth = Peak Bandwidth

Misaligned: two memory transactions

Bandwidth = Peak Bandwidth / 2

Strided access

Bandwidth = Peak Bandwidth / s
for a stride of s

Random access

Bandwidth = Peak Bandwidth / 32

```
__global__ void randomCopy(float* odata, float* idata, int* addr) {  
    int xid = blockIdx.x * blockDim.x + threadIdx.x + offset;  
    odata[xid] = idata[addr[xid]];  
}
```

Summary of main features of device memory

Memory	Location on/off chip	Cached	Access	Scope	Lifetime
Register	On	n/a	R/W	1 thread	Thread
Local	Off	Yes††	R/W	1 thread	Thread
Shared	On	n/a	R/W	All threads in block	Block
Global	Off	†	R/W	All threads + host	Host allocation
Constant	Off	Yes	R	All threads + host	Host allocation
Texture	Off	Yes	R	All threads + host	Host allocation
† Cached in L1 and L2 by default on devices of compute capability 2.x; cached only in L2 by default on devices of higher compute capabilities, though some allow opt-in to caching in L1 as well via compilation flags.					
†† Cached in L1 and L2 by default on devices of compute capability 2.x and 3.x; devices of compute capability 5.x cache locals only in L2.					

Data type size

- The previous discussion covered the most common case where threads read a 4-byte float or int.
- char: 1 byte. float: 4 bytes. double: 8 bytes.
- CUDA has many built-in data types that are shorter/longer:

`char1-4, uchar1-4`

`short1-4, ushort1-4`

`int1-4, uint1-4`

`long1-4, ulong1-4`

`longlong1-2, ulonglong1-2`

`float1-4`

`double1-2`

- Example: `char1, int2, long4, double2`

How to calculate the warp ID

- How do you know if 2 threads are in the same warp?
- Answer: look at the thread ID. Divide by 32: you get the warp ID.
- Thread ID is computed from the 1D, 2d or 3D thread index using:

$$x + yD_x + zD_xD_y$$

```
int tID = threadIdx.x  
        + threadIdx.y * blockDim.x  
        + threadIdx.z * blockDim.x * blockDim.y;  
int warpID = tID/32;
```